

Solution

1. Base cased BERT was used, tuned.
2. 5 classes instead of 3 (B-, I- MOUNTAIN, O),
classes - (B- I- MOUNTAIN, B- I- ELEVATION, O)
3. Trained for 10 epochs, used weighted cross-entropy loss, didn't do any sophisticated hyperparameter tuning.
4. Training dataset consisted of 10k synthetic mount sentences generated using Llama 3.1(check notebooks/data_related), 10k non mount sentences
 - around 2k samples invalidation set of basically same sythetic data that isn't close to real world (if we're not talking about books etc.) (99% accuracy was achieved).
 - Achieved 87% f-score (macro, not weighted) on real world test set from 50 hard samples manually sampled and annotated from subreddit r/alpinism.

How to make better, how i would do it if i start over

- If i started over:

1. Make a list of additional questions for each task (especially CV task).
2. Spend more time for NER trying to manually find and annotate real-world samples.
3. Don't try to learn all needed theory about architectures i'm gonna use.(especially CV DL solutions for matching confused me (papers))

How to make better:

1. Better data (manual sampling (or sophisticated parsing that may introduce bias), annotation)
2. Tuned for NER models
3. Custom losses